


# Supervised Learning: Performance Issues

19 February 2020

# Training Neural Nets

Decisions, decisions...




- How do we...

- Prepare the data?
- Initialise the weights?
- Decide on the architecture?
- Choose activation functions?
- Choose a training algorithm?
- Choose training algorithm parameters?
- Measure model accuracy?
- Measure model complexity?

# Training Neural Nets


Decisions, decisions...



- How do we...
  - Prepare the data?
  - Initialise the weights?
  - Decide on the architecture?
  - Choose activation functions?
  - Choose a training algorithm?
  - Choose training algorithm parameters?
  - Measure model accuracy?
  - Measure model complexity?
- Trial and error

# Training Neural Nets

Decisions, decisions...



- How do we...
  - Prepare the data?
  - Initialise the weights?
  - Decide on the architecture?
  - Choose activation functions?
  - Choose a training algorithm?
  - Choose training algorithm parameters?
  - Measure model accuracy?
  - Measure model complexity?
- Trial and error
- “Whatever I am familiar with!”

# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

## Representation

- Decide on the inputs and the outputs
- Remove obviously irrelevant inputs (names, unique IDs...)

# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

## Representation

- Decide on the inputs and the outputs
- Remove obviously irrelevant inputs (names, unique IDs...)

## Missing Values

Solutions?

# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

## Representation

- Decide on the inputs and the outputs
- Remove obviously irrelevant inputs (names, unique IDs...)

## Missing Values

### Solutions?

- Remove the entire pattern if some data is missing



# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

## Representation

- Decide on the inputs and the outputs
- Remove obviously irrelevant inputs (names, unique IDs...)

## Missing Values

### Solutions?

- Remove the entire pattern if some data is missing
- Replace the missing value with an average for that value over all patterns (continuous) or the most frequently occurring one (discrete)

# Data Preparation

- Correctly pre-processing your data takes you half-way to constructing a good model

## Representation

- Decide on the inputs and the outputs
- Remove obviously irrelevant inputs (names, unique IDs...)

## Missing Values

### Solutions?

- Remove the entire pattern if some data is missing
- Replace the missing value with an average for that value over all patterns (continuous) or the most frequently occurring one (discrete)
- Add an extra input unit to indicate if a parameter is missing

# Data Preparation

## Coding of the Inputs/Outputs

- All inputs/outputs must be numeric

# Data Preparation

## Coding of the Inputs/Outputs

- All inputs/outputs must be numeric
- In case of discrete values/labels, use **binary one-hot** encoding:

# Data Preparation

## Coding of the Inputs/Outputs

- All inputs/outputs must be numeric
- In case of discrete values/labels, use **binary one-hot** encoding:
  - A nominal input that takes  $n$  different values can be coded as  $n$  binary input parameters
  - For each pattern, input corresponding to a specific nominal value is set to 1, the rest are set to 0

# Data Preparation

## Coding of the Inputs/Outputs

- All inputs/outputs must be numeric
- In case of discrete values/labels, use **binary one-hot** encoding:
  - A nominal input that takes  $n$  different values can be coded as  $n$  binary input parameters
  - For each pattern, input corresponding to a specific nominal value is set to 1, the rest are set to 0
- Alternatively, use a single continuous value where every nominal value corresponds to a specific continuous value

# Data Preparation

## Coding of the Inputs/Outputs

- All inputs/outputs must be numeric
- In case of discrete values/labels, use **binary one-hot** encoding:
  - A nominal input that takes  $n$  different values can be coded as  $n$  binary input parameters
  - For each pattern, input corresponding to a specific nominal value is set to 1, the rest are set to 0
- Alternatively, use a single continuous value where every nominal value corresponds to a specific continuous value
  - Discrete characteristic is lost
  - Distance between categories: what does it represent?
  - The representation is more dense (less sparse)

# Data Preparation

## Coding of the Inputs/Outputs


- All inputs/outputs must be numeric
- In case of discrete values/labels, use **binary one-hot** encoding:
  - A nominal input that takes  $n$  different values can be coded as  $n$  binary input parameters
  - For each pattern, input corresponding to a specific nominal value is set to 1, the rest are set to 0
- Alternatively, use a single continuous value where every nominal value corresponds to a specific continuous value
  - Discrete characteristic is lost
  - Distance between categories: what does it represent?
  - The representation is more dense (less sparse)
- Which one is better: **dense or sparse?**



# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques
- Use a robust objective function that is not influenced by outliers




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques
- Use a robust objective function that is not influenced by outliers
  - if  $|E| > \epsilon$ , set  $E$  to a constant value




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques
- Use a robust objective function that is not influenced by outliers
  - if  $|E| > \epsilon$ , set  $E$  to a constant value
  - Problems with this approach?




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques
- Use a robust objective function that is not influenced by outliers
  - if  $|E| > \epsilon$ , set  $E$  to a constant value
  - Problems with this approach?
  - if  $|E| > \epsilon$ , minimise  $|E|$  instead of  $E^2$




# Data Preparation

## Outliers

Outlier patterns produce large errors and divert the search

Solutions:

- Remove outliers using statistical techniques
- Use a robust objective function that is not influenced by outliers
  - if  $|E| > \epsilon$ , set  $E$  to a constant value
  - Problems with this approach?
  - if  $|E| > \epsilon$ , minimise  $|E|$  instead of  $E^2$
  - This approach is called “Huber loss” in statistics



# Data Preparation

## Scaling and Normalization

### Scaling Inputs

- Inputs outside the active domain of the chosen activation function may cause saturation.




# Data Preparation

## Scaling and Normalization

### Scaling Inputs

- Inputs outside the active domain of the chosen activation function may cause saturation.
- What is saturation and why is it bad?




# Data Preparation

## Scaling and Normalization

### Scaling Inputs

- Inputs outside the active domain of the chosen activation function may cause saturation.
- What is saturation and why is it bad?




### Saturation

- Derivatives near asymptotes are close to 0  $\Rightarrow$  slow learning
- A saturated output unit does not indicate the “confidence” level of the NN: all patterns, even the ones not fitted very well by the NN, will be classified with the same “strength”

# Data Preparation

## Linear (Min-Max) Scaling



- Scale inputs/outputs to the necessary range linearly

# Data Preparation


## Linear (Min-Max) Scaling



- Scale inputs/outputs to the necessary range linearly
- $\frac{X-A}{B-A} = \frac{x-a}{b-a}$
- $x = \frac{X-A}{B-A}(b-a) + a$
- $x$  - scaled,  $X$  - unscaled,  $A, B$  - unscaled min and max,  $a, b$  - scaled min and max

# Data Preparation


## Linear (Min-Max) Scaling



- Scale inputs/outputs to the necessary range linearly
- $\frac{X-A}{B-A} = \frac{x-a}{b-a}$
- $x = \frac{X-A}{B-A}(b-a) + a$
- $x$  - scaled,  $X$  - unscaled,  $A, B$  - unscaled min and max,  $a, b$  - scaled min and max
- **Obvious disadvantage:** you need to know the current range of values before you scale them to the desired range

# Data Preparation

## Linear (Min-Max) Scaling



- Scale inputs/outputs to the necessary range linearly
- $\frac{X-A}{B-A} = \frac{x-a}{b-a}$
- $x = \frac{X-A}{B-A}(b-a) + a$
- $x$  - scaled,  $X$  - unscaled,  $A, B$  - unscaled min and max,  $a, b$  - scaled min and max
- **Obvious disadvantage:** you need to know the current range of values before you scale them to the desired range
- What if there is an outlier?

# Data Preparation

## Scaling and Normalization

### Mean Centering: Mean of 0

- Convert the existing distribution to a “gaussian” one
- Average value of variable  $Z_i$  for all  $P$ :  $\bar{Z}_i = \sum_{p=1}^P Z_{i,p} / P$
- Scale:
  - $Z_{i,p}^M = Z_{i,p} - \bar{Z}_i$

# Data Preparation

## Scaling and Normalization

### Mean Centering: Mean of 0

- Convert the existing distribution to a “gaussian” one
- Average value of variable  $Z_i$  for all  $P$ :  $\bar{Z}_i = \sum_{p=1}^P Z_{i,p} / P$
- Scale:
  - $Z_{i,p}^M = Z_{i,p} - \bar{Z}_i$

### Variance Scaling: Variance of 1

- Let  $\sigma_{Z_i}$  be the standard deviations of  $Z_{i,p}$ . Then:
  - $Z_{i,p}^V = \frac{Z_{i,p}}{\sigma_{Z_i}}$




# Data Preparation

Combine mean centering and variance scaling

## Z-score (“standard score”) normalization

- Combine mean centering and variance scaling to normalize the data:

$$Z_{i,p}^{MV} = \frac{Z_{i,p} - \bar{Z}_i}{\sigma_{Z_i}}$$




# Data Preparation

## Scaling and Normalization

### Scaling Outputs

- Outputs outside of the activation function range will be unreachable
- NN will always produce a large error




# Data Preparation

## Scaling and Normalization

### Scaling Outputs

- Outputs outside of the activation function range will be unreachable
- NN will always produce a large error



### Outputs for bounded functions


- Make outputs reachable: scale to  $[0, 1]$  for Sigmoid,  $[-1, 1]$  for TanH, etc.

# Data Preparation

## Scaling and Normalization

### Scaling Outputs

- Outputs outside of the activation function range will be unreachable
- NN will always produce a large error



### Outputs for bounded functions

- Make outputs reachable: scale to  $[0, 1]$  for Sigmoid,  $[-1, 1]$  for TanH, etc.
- Trick of the trade: scale to  $[0.1, 0.9]$  (Sigmoid) and  $[-0.9, 0.9]$  (TanH) instead (why?)

# Data Preparation

Do you have enough data?

- Many diverse examples are better than a few similar examples

# Data Preparation

Do you have enough data?

- Many diverse examples are better than a few similar examples
- The existing data set can always be artificially expanded
  - Add **noise** sampled from a normal distribution with a small variance and zero mean to the existing patterns

# Data Preparation

Do you have enough data?

- Many diverse examples are better than a few similar examples
- The existing data set can always be artificially expanded
  - Add **noise** sampled from a normal distribution with a small variance and zero mean to the existing patterns
  - Image data: rotate, distort, zoom in and out, etc.

# Data Preparation

Do you have enough data?


- Many diverse examples are better than a few similar examples
- The existing data set can always be artificially expanded
  - Add **noise** sampled from a normal distribution with a small variance and zero mean to the existing patterns
  - Image data: rotate, distort, zoom in and out, etc.
- Adding random noise at every epoch helps prevent overfitting - it becomes harder to memorise the exact patterns, so the NN learns the bigger picture instead



# Data Preparation

Do you have enough data?

- Many diverse examples are better than a few similar examples
- The existing data set can always be artificially expanded
  - Add **noise** sampled from a normal distribution with a small variance and zero mean to the existing patterns
  - Image data: rotate, distort, zoom in and out, etc.
- Adding random noise at every epoch helps prevent overfitting - it becomes harder to memorise the exact patterns, so the NN learns the bigger picture instead
- Deep learning: first record-breaking MNIST-recognizing NNs modified the data set by artificial expansion (rotations, distortions, etc.)
- <http://yann.lecun.com/exdb/mnist/>



# Weight Initialisation

## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal

# Weight Initialisation

## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal
- Population-based algorithms such as PSO and GA use a population of points: not distributing them adequately can cause stagnation/premature convergence

# Weight Initialisation

## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal
- Population-based algorithms such as PSO and GA use a population of points: not distributing them adequately can cause stagnation/premature convergence
- Can you set all weights to zero?

# Weight Initialisation

## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal
- Population-based algorithms such as PSO and GA use a population of points: not distributing them adequately can cause stagnation/premature convergence
- Can you set all weights to zero?
- Can you set weights to random numbers?

# Weight Initialisation

## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal
- Population-based algorithms such as PSO and GA use a population of points: not distributing them adequately can cause stagnation/premature convergence
- Can you set all weights to zero?
- Can you set weights to random numbers?
- Random weights in a small range around 0:
  - Small net input signals
  - Midrange values produced by activation functions
  - I.e. no instant saturation => better learning potential

# Weight Initialisation


## Choosing the starting point

- Gradient descent starts on a single point: choosing a “bad” starting point can be fatal
- Population-based algorithms such as PSO and GA use a population of points: not distributing them adequately can cause stagnation/premature convergence
- Can you set all weights to zero?
- Can you set weights to random numbers?
- Random weights in a small range around 0:
  - Small net input signals
  - Midrange values produced by activation functions
  - I.e. no instant saturation => better learning potential

# Weight Initialisation

## Choosing the starting point

- Gaussian distribution can be used instead of uniform to sample the weights:






# Weight Initialisation

## Choosing the starting point

- Gaussian distribution can be used instead of uniform to sample the weights:



- **Wessels and Barnard**: choose random weights in range  $\left[ \frac{-1}{\sqrt{fanin}}, \frac{1}{\sqrt{fanin}} \right]$ , where *fanin* is the number of incoming connections for the given unit.
  - The larger the architecture, the smaller interval will be used

# Weight Initialisation

## Choosing the starting point

- **Xavier and Glorot (2010)**: sample random weights from a distribution with  $\mu = 0$  and  $\sigma = \frac{\sqrt{6}}{\sqrt{fanin+fanout}}$ .
  - Derived to normalise (maintain variance of) the backpropagated gradients

# Weight Initialisation

## Choosing the starting point

- **Xavier and Glorot (2010)**: sample random weights from a distribution with  $\mu = 0$  and  $\sigma = \frac{\sqrt{6}}{\sqrt{fanin+fanout}}$ .
  - Derived to normalise (maintain variance of) the backpropagated gradients
- **He Normal (He et al., 2015)** [for rectified activations]:
  - 1 Choose weights randomly from a standard normal distribution.
  - 2 Multiply each weight by  $\frac{\sqrt{2}}{\sqrt{fanin}}$ .
  - 3 Bias weights are initialized to zero.

# Weight Initialisation

## Choosing the starting point

- **Xavier and Glorot (2010)**: sample random weights from a distribution with  $\mu = 0$  and  $\sigma = \frac{\sqrt{6}}{\sqrt{fanin+fanout}}$ .
  - Derived to normalise (maintain variance of) the backpropagated gradients
- **He Normal (He et al., 2015)** [for rectified activations]:
  - 1 Choose weights randomly from a standard normal distribution.
  - 2 Multiply each weight by  $\frac{\sqrt{2}}{\sqrt{fanin}}$ .
  - 3 Bias weights are initialized to zero.
- **Are smaller weights always better?**

# Data set

## Training, Generalisation, Validation

- Supervised training: data patterns with known target values are presented to the NN
- Data set has to be subdivided into three parts:

### Training set

Data that the training algorithm will use to iteratively adjust the weights

# Data set

## Training, Generalisation, Validation

- Supervised training: data patterns with known target values are presented to the NN
- Data set has to be subdivided into three parts:

### Training set

Data that the training algorithm will use to iteratively adjust the weights

### Generalisation/Test set

Data that will be used to calculate the generalisation error during hyperparameter optimization

# Data set

## Training, Generalisation, Validation

- Supervised training: data patterns with known target values are presented to the NN
- Data set has to be subdivided into three parts:

### Training set

Data that the training algorithm will use to iteratively adjust the weights

### Generalisation/Test set

Data that will be used to calculate the generalisation error during hyperparameter optimization

### Validation set

Data that will be used to calculate the generalisation error of the optimized model

# Data set

Training, Generalisation, Validation

## Separating the data

Using training/optimization data for testing is not fair:  
generalisation error will not be an objective measure



# Data set

Training, Generalisation, Validation

## Separating the data

Using training/optimization data for testing is not fair:  
generalisation error will not be an objective measure

## Data normalisation

How would you scale the data in training, testing, validation subsets?

# Data Scaling

Which one is correct?

# Data Scaling

Which one is correct?

- 1 Scale the whole data set before splitting:

$$T_i = \frac{T_i - \bar{T}_{all}}{\sigma_{all}}$$

# Data Scaling

Which one is correct?

- 1 Scale the whole data set before splitting:

$$T_i = \frac{T_i - \bar{T}_{all}}{\sigma_{all}}$$

- 2 Scale training and testing sets separately:

$$T_{i,train} = \frac{T_{i,train} - \bar{T}_{train}}{\sigma_{train}}, T_{i,test} = \frac{T_{i,test} - \bar{T}_{test}}{\sigma_{test}}$$

# Data Scaling

Which one is correct?

- 1 Scale the whole data set before splitting:

$$T_i = \frac{T_i - \bar{T}_{all}}{\sigma_{all}}$$

- 2 Scale training and testing sets separately:

$$T_{i,train} = \frac{T_{i,train} - \bar{T}_{train}}{\sigma_{train}}, T_{i,test} = \frac{T_{i,test} - \bar{T}_{test}}{\sigma_{test}}$$

- 3 Apply training data transformation across the board:

$$T_{i,train} = \frac{T_{i,train} - \bar{T}_{train}}{\sigma_{train}}, T_{i,test} = \frac{T_{i,test} - \bar{T}_{train}}{\sigma_{train}}$$

# Data Scaling

Which one is correct?

- ① Scale the whole data set before splitting:

$$T_i = \frac{T_i - \bar{T}_{all}}{\sigma_{all}}$$

- ② Scale training and testing sets separately:

$$T_{i,train} = \frac{T_{i,train} - \bar{T}_{train}}{\sigma_{train}}, T_{i,test} = \frac{T_{i,test} - \bar{T}_{test}}{\sigma_{test}}$$

- ③ Apply training data transformation across the board:

$$T_{i,train} = \frac{T_{i,train} - \bar{T}_{train}}{\sigma_{train}}, T_{i,test} = \frac{T_{i,test} - \bar{T}_{train}}{\sigma_{train}}$$

**Option (3) is the correct approach.**

# Data set

## Training, Generalisation, Validation

- Training set + generalisation set make up the training set for the optimized model
- How do we subdivide the data set?

### Training set

Over 50%, up to 90% of the entire data set

# Data set

## Training, Generalisation, Validation

- Training set + generalisation set make up the training set for the optimized model
- How do we subdivide the data set?

### Training set

Over 50%, up to 90% of the entire data set

### Generalisation set

About 10% to 30% of the training set



# Data set

## Training, Generalisation, Validation

- Training set + generalisation set make up the training set for the optimized model
- How do we subdivide the data set?

### Training set

Over 50%, up to 90% of the entire data set

### Generalisation set

About 10% to 30% of the training set

### Validation set

About 10% to 30% of the data set. Should never be used for training.

# Training the NN

Stochastic, batch, mini-batch

## Stochastic (on-line) training

- Update weights after each pattern
- Backprop each gradient
- Bound to fluctuate: is a single pattern representative?

# Training the NN

Stochastic, batch, mini-batch

## Stochastic (on-line) training

- Update weights after each pattern
- Backprop each gradient
- Bound to fluctuate: is a single pattern representative?

## Batch training

- Update weights after all patterns were considered
- Backprop average gradient
- Stable, but computationally heavy

# Training the NN

Stochastic, batch, mini-batch

## Stochastic (on-line) training

- Update weights after each pattern
- Backprop each gradient
- Bound to fluctuate: is a single pattern representative?

## Batch training

- Update weights after all patterns were considered
- Backprop average gradient
- Stable, but computationally heavy

## Mini-Batch training: Best of both

- Calculate average gradient over a subset of patterns
- Very popular in deep learning

# Training the NN

Stochastic, batch, mini-batch

## Shuffling the data

- Training set must be shuffled at every epoch (iteration)

# Training the NN

Stochastic, batch, mini-batch

## Shuffling the data

- Training set must be shuffled at every epoch (iteration)
- If patterns are presented in the same order, the NN may infer the order of patterns and learn it

# Training the NN

Stochastic, batch, mini-batch

## Shuffling the data

- Training set must be shuffled at every epoch (iteration)
- If patterns are presented in the same order, the NN may infer the order of patterns and learn it
- If subsets of the data set are used, the subsets may not be representative unless the data set is shuffled

# Training the NN

Stochastic, batch, mini-batch

## Shuffling the data

- Training set must be shuffled at every epoch (iteration)
- If patterns are presented in the same order, the NN may infer the order of patterns and learn it
- If subsets of the data set are used, the subsets may not be representative unless the data set is shuffled
- How does this apply to stochastic, batch, mini-batch training?